

Aido Data

DFoly

June 30, 2016

We want to predict whether or not there will be fraud based on the variables in the Data.

Initial Data Exploration

```
Data <- read.csv("C:/Users/dfoley/Dropbox/Aido/Data_Presentation/Fraud_data.csv")
summary(Data)
```

```
##          FRAUD          PPSN    SCHEME_CODE LOCATION_CODE
## Min.      :0.000    0001135J:    1    F1:1113      D1       : 28
## 1st Qu.:0.000    0003334Z:    1    G1: 99       S1       : 26
## Median :1.000    0006622M:    1    H1: 494      A3       : 24
## Mean    :0.509    0009177V:    1    H2: 294      G4       : 24
## 3rd Qu.:1.000    0024727R:    1           V3       : 23
## Max.    :1.000    0028683D:    1           E3       : 22
##          (Other) :1994          (Other):1853
##          PAY_TYPE      REG_DATE      START_DATE      BIRTH_DATE
## Basic      :1102    07/09/2008:    4    10/04/2001:    4    02/02/1995:    3
## Exceptional: 253    15/07/2009:    4    12/05/2012:    4    12/04/1977:    3
## Supplement : 472    23/04/2012:    4    22/03/2011:    4    15/01/1986:    3
## Urgent     : 173    01/06/2007:    3    22/05/2012:    4    24/06/1973:    3
##           02/02/2008:    3    01/11/2012:    3    25/03/1983:    3
##           02/03/2012:    3    02/05/2009:    3    28/06/1982:    3
##           (Other)   :1979    (Other)   :1978    (Other)   :1982
##          AGE_YEARS    OCC_CODE    MEANS_FROM_EMP      DELAY
## Min.      :18.00    M :779    Min.      : 100    Min.      : 0.00
## 1st Qu.:29.00    P  : 98    1st Qu.:10700    1st Qu.:14.00
## Median :37.00    S  :435    Median :21150    Median :21.00
## Mean     :38.67    SS :687    Mean   :21595    Mean   :21.86
## 3rd Qu.:47.00    NA's: 1    3rd Qu.:31400    3rd Qu.:29.00
## Max.     :79.00           Max.   :57800    Max.   :60.00
##
##          WEEKLY_RATE    PAID_TO_DATE      SEX      MARITAL_STATUS
## Min.      : 50.0    Min.      : 4160    Female: 974    D:666
## 1st Qu.:120.0    1st Qu.: 29160    Male  :1026    M:446
## Median :160.0    Median : 49390           S:888
## Mean     :154.3    Mean   : 59593
## 3rd Qu.:190.0    3rd Qu.: 81600
## Max.     :220.0    Max.   :220290
## NA's      :1      NA's      :1
##          NO_DEP      FUEL_AMOUNT      APPOINTMENT_STATUS    BOOK_NUMBER
## Min.      : 0.00    Min.      : 40.00    Due          :746    Min.      :10017
## 1st Qu.: 2.00    1st Qu.: 70.00    Missed       :391    1st Qu.:19951
## Median : 3.00    Median : 80.00    Up-to-Date:863    Median :29610
```

```
## Mean : 3.36 Mean : 78.06 Mean :29721
## 3rd Qu.: 4.00 3rd Qu.: 90.00 3rd Qu.:39568
## Max. :10.00 Max. :140.00 Max. :49995
##
## CARER_ID PHOTO_VERIFIED OPEN_CREDIT_LINES DRIVING_STATUS
## Min. : 0 Min. :0.000 Min. : 0.000 Active :1402
## 1st Qu.: 0 1st Qu.:0.000 1st Qu.: 4.000 None : 518
## Median : 0 Median :1.000 Median : 5.000 Suspended: 80
## Mean : 1717 Mean :0.697 Mean : 5.082
## 3rd Qu.: 0 3rd Qu.:1.000 3rd Qu.: 6.000
## Max. :49966 Max. :1.000 Max. :11.000
##
## FREE_TRAVEL EMPLOYER_NO IBAN MED_ASSESSMENT
## Min. :0.000 Min. : 18833 AAAA080879195849: 1 Min. :1.00
## 1st Qu.:0.000 1st Qu.:25438583 AAAA294430715538: 1 1st Qu.:1.00
## Median :0.000 Median :50508358 AAAA504389685081: 1 Median :2.00
## Mean :0.219 Mean :50085185 AAAB021425204220: 1 Mean :2.31
## 3rd Qu.:0.000 3rd Qu.:74127544 AAAB109040882998: 1 3rd Qu.:3.00
## Max. :1.000 Max. :99989585 AAAB190156605237: 1 Max. :4.00
## (Other) :1994
## MED_CONDITION_SATISFIED MORT_OUTSTANDING PAY_LOCATION_CODE PREF_LANG
## Min. :0.0000 Min. : 0 D1 : 28 Min. :1.00
## 1st Qu.:0.0000 1st Qu.: 0 V3 : 26 1st Qu.:1.00
## Median :0.0000 Median : 0 S1 : 25 Median :1.00
## Mean :0.1605 Mean : 58189 A3 : 23 Mean :1.03
## 3rd Qu.:0.0000 3rd Qu.:137688 E3 : 23 3rd Qu.:1.00
## Max. :1.0000 Max. :256947 Q2 : 23 Max. :2.00
## (Other):1852
## PHONE_NO PRIORITY_CLAIM
## Min. :3.54e+11 Min. :0.000
## 1st Qu.:3.54e+11 1st Qu.:0.000
## Median :3.54e+11 Median :0.000
## Mean :3.54e+11 Mean :0.096
## 3rd Qu.:3.54e+11 3rd Qu.:0.000
## Max. :3.54e+11 Max. :1.000
##
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.1.3
```

Lets start off with an initial exploration of the data to see what variables may impact fraud

```
cor(Data[,unlist(lapply(Data, is.numeric))], Data$FRAUD)
```

```
## Warning in cor(Data[, unlist(lapply(Data, is.numeric))], Data$FRAUD): the
## standard deviation is zero
```

```
## [1]
## FRAUD 1.0000000000
## AGE_YEARS -0.022102019
```

```
## MEANS_FROM_EMP      -0.691814409
## DELAY               0.557881319
## WEEKLY_RATE         NA
## PAID_TO_DATE        NA
## NO_DEP              0.398714291
## FUEL_AMOUNT         0.213654006
## BOOK_NUMBER         0.035143006
## CARER_ID            -0.018820769
## PHOTO_VERIFIED      0.003164441
## OPEN_CREDIT_LINES   0.012542585
## FREE_TRAVEL         -0.011951586
## EMPLOYER_NO         -0.023686280
## MED_ASSESSMENT      0.149375417
## MED_CONDITION_SATISFIED -0.259908686
## MORT_OUTSTANDING    -0.001932028
## PREF_LANG           -0.006094250
## PHONE_NO            NA
## PRIORITY_CLAIM      -0.053397883
```

A few variables look to be correlated with FRAUD

In particular:

MEANS_FROM_EMP

DELAY

NO_DEP

FUEL_AMOUNT

MED_ASSESSMENT

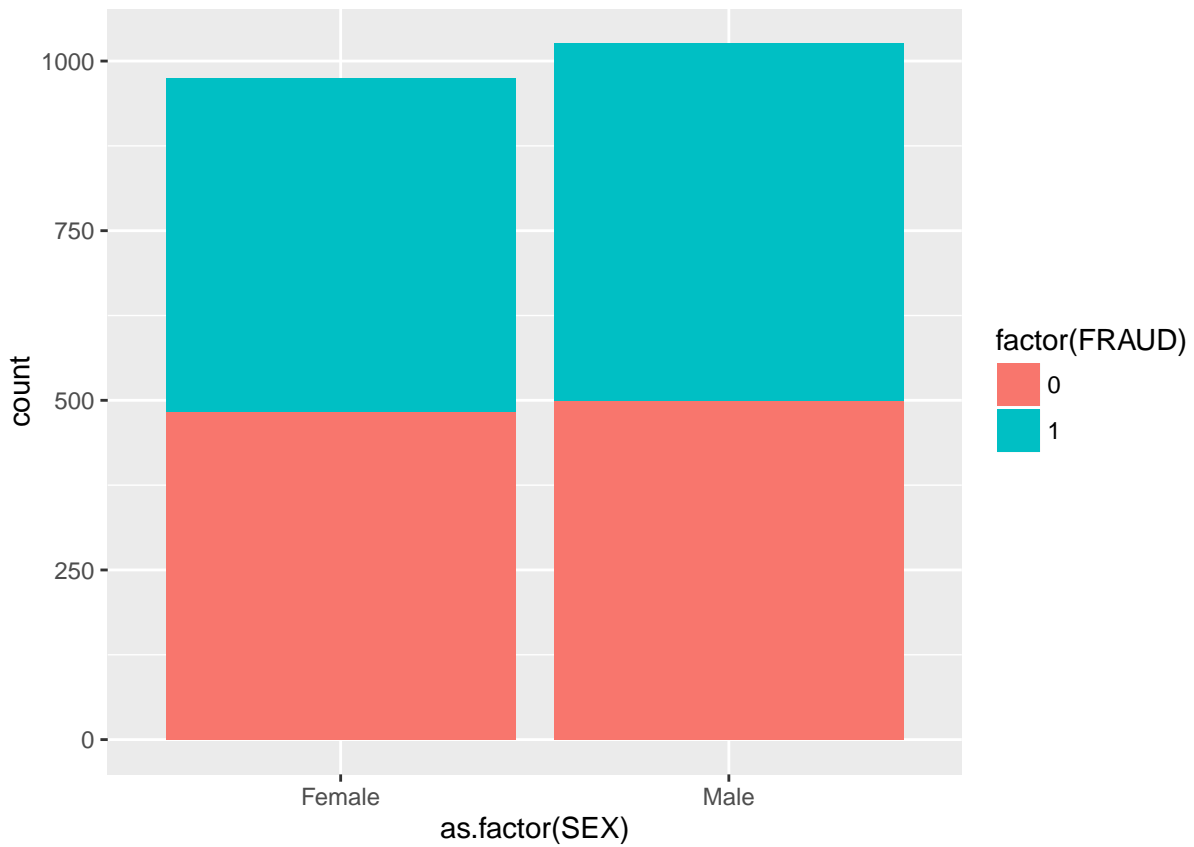
MED_CONDITION_SATISFIED

No other variables appear to have any strong correlation

```
table(Data$FRAUD, Data$SEX)
```


```
##
##      Female Male
## 0      483  499
## 1      491  527
```

```
ggplot(Data, aes(x = as.factor(SEX), fill=factor(FRAUD)))+geom_bar()
```




About 49% of Males and Females commit fraud

```
ggplot(Data, aes(x = AGE_YEARS, fill = factor(FRAUD))) + geom_bar()
```




```
ggplot(Data, aes(x = WEEKLY_RATE, fill = as.factor(FRAUD))) +  
  geom_bar(stat = 'count', position = 'dodge')
```

```
## Warning: Removed 1 rows containing non-finite values (stat_count).
```




The level of fraud doesnt seem to be particularly more prevalent in any age group or at any WEEKLY_RATE.

```
ggplot(Data, aes(x = factor(FRAUD), y = MEANS_FROM_EMP)) +  
  geom_boxplot() + geom_hline(aes(yintercept = mean(MEANS_FROM_EMP)),  
                                linetype = 'dashed', lwd = 2)
```



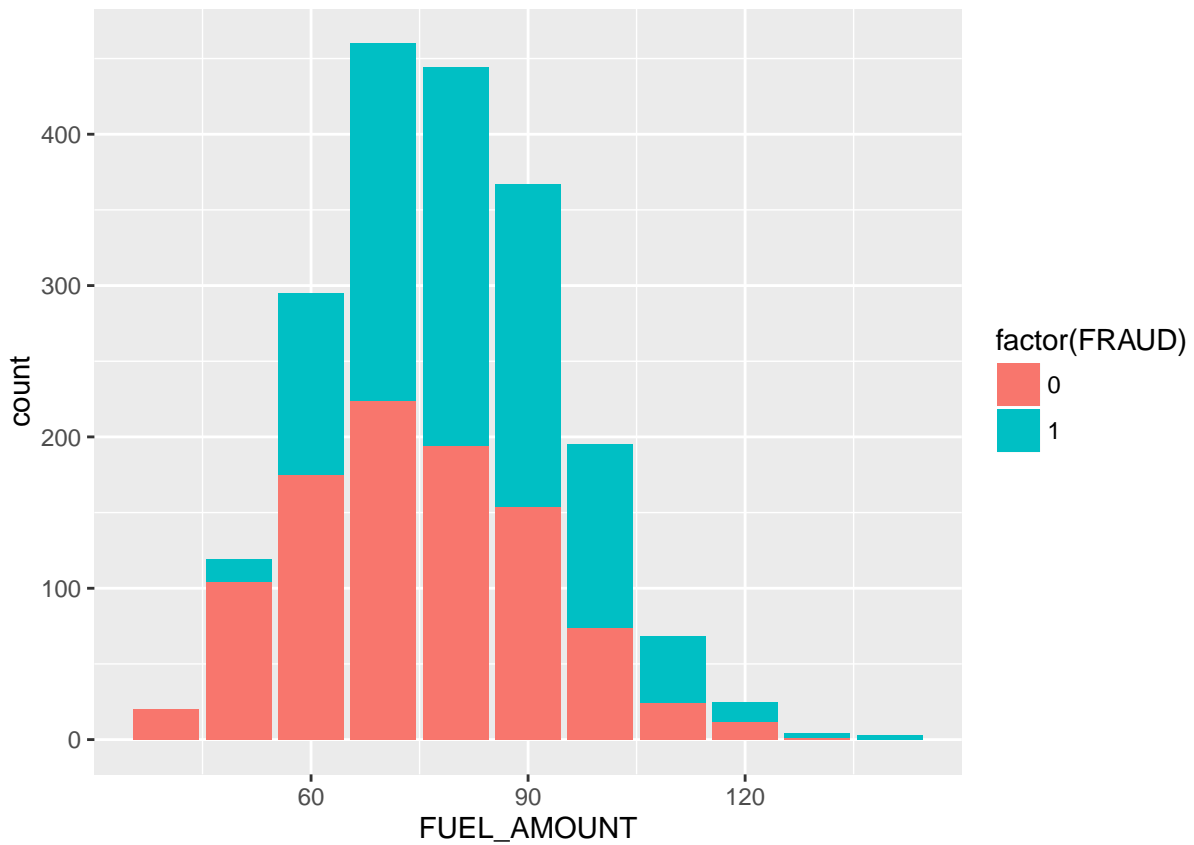
clearly more likely to commit fraud if below the mean value
of MEAN_FROM_EMP

```
ggplot(Data, aes(x = DELAY, fill = factor(FRAUD))) + geom_bar()
```



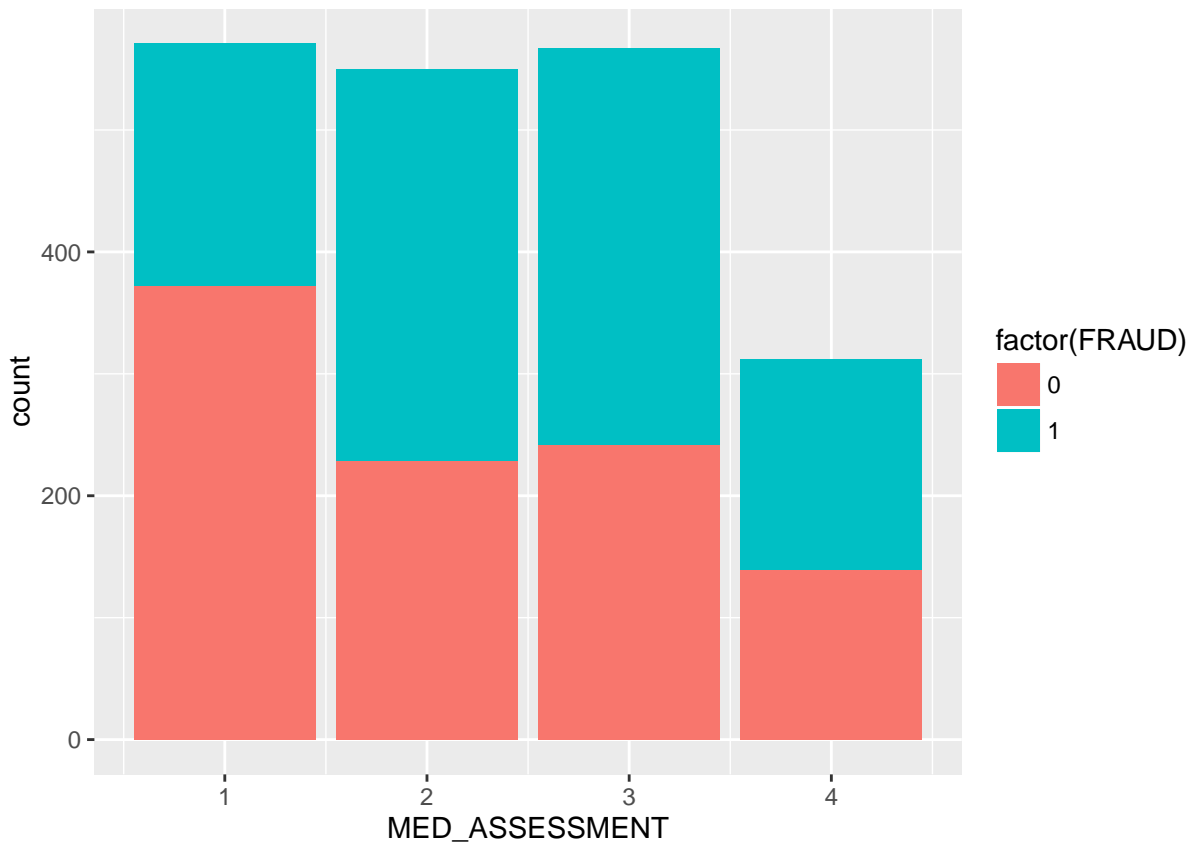
As delay increases there are increasingly higher proportion of FRAUDS

```
ggplot(Data, aes(x = FUEL_AMOUNT, fill = factor(FRAUD))) + geom_bar()
```


There is also perhaps some relevant information we can use here especially at higher fuel levels.

```
ggplot(Data, aes(x = MED_ASSESSMENT, fill = factor(FRAUD))) + geom_bar()
```




If you score one in medical assesment there is a lower chance you will commit FRAUD.

```
ggplot(Data, aes(x = factor(MED_CONDITION_SATISFIED), fill = factor(FRAUD))) + geom_bar()
```



Assuming 1 is having satisfied medical condition, you are less likely to commit fraud.

```
ggplot(Data, aes(x = factor(MED_CONDITION_SATISFIED), fill = factor(FRAUD))) + geom_bar()
```



Data Cleaning

Check pattern of missing values using Amelia package

```
library(Amelia)
```

```
## Warning: package 'Amelia' was built under R version 3.1.3
```


```
## Loading required package: Rcpp
```

```
## Warning: package 'Rcpp' was built under R version 3.1.3
```

```
## ##
## ## Amelia II: Multiple Imputation
## ## (Version 1.7.4, built: 2015-12-05)
## ## Copyright (C) 2005-2016 James Honaker, Gary King and Matthew Blackwell
## ## Refer to http://gking.harvard.edu/amelia/ for more information
## ##
```

```
missmap(Data, main="Fraud Data missing values",
         col=c("blue", "red"), legend=FALSE)
```

Fraud Data missing values



It doesn't look like there is any data missing from the variables we are going to be working with
 NO_DEP and DELAY have zero values so we should inspect them to see if they are errors or not

```
summary(Data$MEANS_FROM_EMP)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      100  10700   21150   21590   31400   57800
```

```
Data$MEANS_FROM_EMP <- scale(Data$MEANS_FROM_EMP)
```

They don't seem to be errors, however all values of zero are not frauds
 We have also normalised the MEANS_FROM_EMP for any algorithms we use

Try Logistic Regression

Split into training and test sets (90/10)

```
Data$MED_ASSESSMENT <- as.factor(Data$MED_ASSESSMENT)
Data$MED_CONDITION_SATISFIED <- as.factor(Data$MED_CONDITION_SATISFIED)
Data$NO_DEP <- as.factor(Data$NO_DEP)
```

```
Data$MEANS_NEW <- ifelse(Data$MEANS_FROM_EMP <= mean(Data$MEANS_FROM_EMP),1,0)
brks <- c(0,20,40,60)
Data$DELAY_RANGE <- cut(Data$DELAY, breaks = brks, include.lowest = T)
summary(Data$DELAY_RANGE )
```

```
## [0,20] (20,40] (40,60]
##      969      919      112
```

```
brks2 <- c(0,60,100,140)
Data$FUEL_NEW <- cut(Data$FUEL_AMOUNT, breaks = brks2, include.lowest = T)
summary(Data$FUEL_NEW)
```

```
## [0,60] (60,100] (100,140]
##      434      1466      100
```

```
Data$NO_DEP2 <- ifelse(as.numeric(Data$NO_DEP) > 4,1,0)
Data$NO_DEP2 <- as.factor(Data$NO_DEP2)
```

```
Data$gp <- runif(dim(Data)[1])
trainingSet <- subset(Data, Data$gp > 0.1)
testSet <- subset(Data, Data$gp <= 0.1)
```

```
reg1 <- glm(FRAUD ~ MEANS_FROM_EMP + DELAY + NO_DEP +
            FUEL_AMOUNT + MED_ASSESSMENT + MED_CONDITION_SATISFIED,
            data = trainingSet, family = binomial(link = 'logit'))
summary(reg1)
```

```
##
## Call:
## glm(formula = FRAUD ~ MEANS_FROM_EMP + DELAY + NO_DEP + FUEL_AMOUNT +
##      MED_ASSESSMENT + MED_CONDITION_SATISFIED, family = binomial(link = "logit"),
##      data = trainingSet)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.96907  -0.23256   0.02181   0.25794   3.09717
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -16.98086   598.48538  -0.028   0.9774
## MEANS_FROM_EMP    -2.38553    0.13989 -17.053 < 2e-16 ***
## DELAY              0.15008    0.01163  12.906 < 2e-16 ***
## NO_DEP1           15.43118   598.48534   0.026   0.9794
## NO_DEP2           18.24755   598.48523   0.030   0.9757
## NO_DEP3           19.62547   598.48529   0.033   0.9738
## NO_DEP4           20.67963   598.48539   0.035   0.9724
## NO_DEP5           21.98400   598.48551   0.037   0.9707
## NO_DEP6           22.97988   598.48570   0.038   0.9694
## NO_DEP7           24.32543   598.48664   0.041   0.9676
## NO_DEP8           38.59980  1333.36553   0.029   0.9769
```

```
## NO_DEP9                45.02335 4001.19338 0.011 0.9910
## NO_DEP10               37.56084 4001.19339 0.009 0.9925
## FUEL_AMOUNT            -0.08028 0.01149 -6.990 2.75e-12 ***
## MED_ASSESSMENT2        0.35933 0.26482 1.357 0.1748
## MED_ASSESSMENT3        0.60503 0.26378 2.294 0.0218 *
## MED_ASSESSMENT4        0.43600 0.30252 1.441 0.1495
## MED_CONDITION_SATISFIED1 -1.31575 0.30904 -4.257 2.07e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 2488.80 on 1795 degrees of freedom
## Residual deviance: 835.23 on 1778 degrees of freedom
## AIC: 871.23
##
## Number of Fisher Scoring iterations: 16
```

There seems to be significance in all but one variable

Test Fit

```
fit1 <- predict(reg1, newdata = testSet, type = 'response')
fit1 <- ifelse(fit1 > 0.5,1,0)
misClasificError1 <- mean(fit1 != testSet$FRAUD)
print(paste('Accuracy',1-misClasificError1))
```

```
## [1] "Accuracy 0.901960784313726"
```

Logistic Part 2

Try and increase performance by using logit on redfined variables

```
reg2 <- glm(FRAUD ~ MEANS_NEW + DELAY_RANGE + NO_DEP2 +
            FUEL_NEW + MED_ASSESSMENT + MED_CONDITION_SATISFIED,
            data = trainingSet, family = binomial(link = 'logit'))
summary(reg2)
```

```
##
## Call:
## glm(formula = FRAUD ~ MEANS_NEW + DELAY_RANGE + NO_DEP2 + FUEL_NEW +
## MED_ASSESSMENT + MED_CONDITION_SATISFIED, family = binomial(link = "logit"),
## data = trainingSet)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.66448  -0.39811   0.00008   0.46829   2.71299
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.6546    0.2713 -13.473  < 2e-16 ***
```

```
## MEANS_NEW          3.0138      0.1530  19.699 < 2e-16 ***
## DELAY_RANGE(20,40] 1.9520      0.1507  12.949 < 2e-16 ***
## DELAY_RANGE(40,60] 18.1385    337.3166   0.054  0.95712
## NO_DEP21          1.0501      0.1703   6.166 7.00e-10 ***
## FUEL_NEW(60,100]   0.4862      0.2010   2.419  0.01557 *
## FUEL_NEW(100,140]  0.2435      0.3873   0.629  0.52963
## MED_ASSESSMENT2    0.5724      0.2187   2.617  0.00886 **
## MED_ASSESSMENT3    0.6731      0.2194   3.067  0.00216 **
## MED_ASSESSMENT4    0.3578      0.2517   1.422  0.15510
## MED_CONDITION_SATISFIED1 -1.1811    0.2519  -4.689 2.75e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 2488.8  on 1795  degrees of freedom
## Residual deviance: 1224.4  on 1785  degrees of freedom
## AIC: 1246.4
##
## Number of Fisher Scoring iterations: 16
```

Test Fit

```
fit <- predict(reg2, newdata = testSet, type = 'response')
fit <- ifelse(fit> 0.5,1,0)

misClasificError <- mean(fit != testSet$FRAUD)
print(paste('Accuracy',1-misClasificError))
```

```
## [1] "Accuracy 0.838235294117647"
```

Random Forest

Next we wil try a randomforest wich tend to perform well with this kind of problem

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.1.3
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##    margin
```

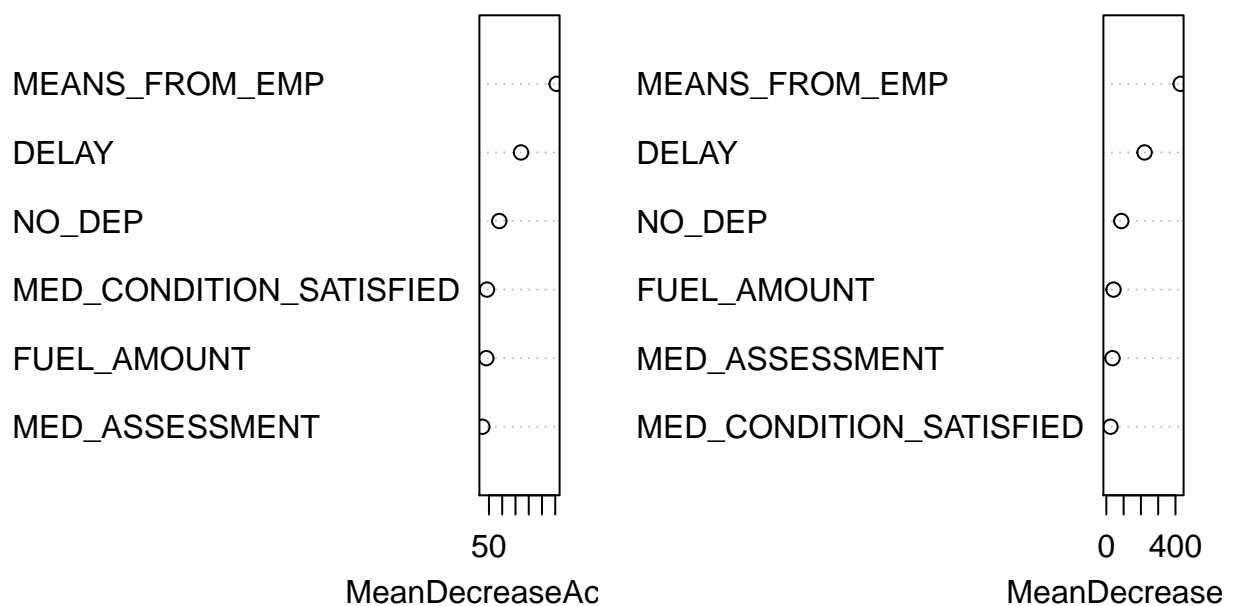


```

set.seed(111)
rf2 <- randomForest(as.factor(FRAUD) ~ MEANS_FROM_EMP + DELAY + NO_DEP +
  FUEL_AMOUNT + MED_ASSESSMENT + MED_CONDITION_SATISFIED,
  data = trainingSet,
  importance = TRUE,
  ntree = 2000)
varImpPlot(rf2)

```

rf2



```

fit3 <- predict(rf2, newdata = testSet, type = 'response')
misClasificError3 <- mean(fit3 != testSet$FRAUD)
print(paste('Accuracy', 1-misClasificError3))

```

```
## [1] "Accuracy 0.901960784313726"
```

Accuracy was reduced slightly

Finally we perform a conditional random forest from the party package

```
library(party)
```

```
## Warning: package 'party' was built under R version 3.1.3
```

```
## Loading required package: grid
```

```
## Loading required package: mvtnorm

## Warning: package 'mvtnorm' was built under R version 3.1.3

## Loading required package: modeltools

## Warning: package 'modeltools' was built under R version 3.1.3

## Loading required package: stats4

## Loading required package: strucchange

## Warning: package 'strucchange' was built under R version 3.1.3

## Loading required package: zoo

## Warning: package 'zoo' was built under R version 3.1.3

##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric

## Loading required package: sandwich

## Warning: package 'sandwich' was built under R version 3.1.3
```

```
set.seed(333)
rf1 <- cforest(as.factor(FRAUD) ~ MEANS_FROM_EMP + DELAY + NO_DEP +
              FUEL_AMOUNT + MED_ASSESSMENT + MED_CONDITION_SATISFIED,
              data = trainingSet,
              controls=cforest_unbiased(ntree=2000, mtry=3))

fit4 <- predict(rf1, newdata = testSet, type = 'response')
misClasificError4 <- mean(fit4 != testSet$FRAUD)
print(paste('Accuracy', 1-misClasificError4))
```

```
## [1] "Accuracy 0.887254901960784"
```

It looks as though our original Logistic Regression is actually the most accurate on the tests set. We can perform Cross validation to confirm.

```
formula = as.factor(FRAUD) ~ MEANS_NEW + DELAY_RANGE +
          NO_DEP2 + FUEL_NEW + MED_ASSESSMENT +
          MED_CONDITION_SATISFIED

library(caret)
```

```
## Warning: package 'caret' was built under R version 3.1.3
```

```
## Loading required package: lattice
```

```
train_control <- trainControl(method = 'cv', number = 10)
modelLog = train(formula, data=trainingSet, method="glm", family=binomial,
                  trControl=train_control)
print(modelLog)
```

```
## Generalized Linear Model
```

```
##
```

```
## 1796 samples
```

```
## 38 predictor
```

```
## 2 classes: '0', '1'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 1616, 1616, 1616, 1617, 1617, 1616, ...
```

```
## Resampling results
```

```
##
```

```
## Accuracy Kappa Accuracy SD Kappa SD
## 0.8529857 0.7056808 0.02322506 0.04658199
```

```
##
```

```
##
```

Try on RandomForest

```
train_control2 <- trainControl(method = 'cv', number = 10)
modelLog2 = train(formula, data=trainingSet, method="rf",
                  trControl=train_control)
print(modelLog2)
```

```
## Random Forest
```

```
##
```

```
## 1796 samples
```

```
## 38 predictor
```

```
## 2 classes: '0', '1'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 1617, 1616, 1617, 1616, 1616, 1617, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
## mtry Accuracy Kappa Accuracy SD Kappa SD
## 2 0.8457604 0.6910251 0.02439651 0.04908086
## 6 0.8418560 0.6836242 0.02988816 0.05984692
## 10 0.8412973 0.6825154 0.03136126 0.06276969
```

```
##
```

```
## Accuracy was used to select the optimal model using the largest value.
```

```
## The final value used for the model was mtry = 2.
```